

Survival Selection Grows a Private Kin Code, Not a Public Language

Chunxuan Yang

Sogang University, Seoul, Republic of Korea
yang2023@sogang.ac.kr yangchunxuan1@gmail.com

Abstract

Can a decodable signalling code *grow* from random weights under survival selection alone, and is what grows public or private to kin? In a minimal hand-coded 2-D lethal-vent world, tiny tied code-table agents evolve a decodable channel. Against an architecture-matched random-fitness control (differing only in the reproduction-fitness signal), selection *sharpens* rather than originates it: clonal descent alone reaches mutual intelligibility (MII) 0.1174 ($\sim 73\times$ the 1/625 floor), and survival selection adds a paired win-margin of +0.0548 (95% CI [0.0234, 0.0849], $n=48$) at gen 28, a margin a post-hoc probe shows is a *transient acceleration*. The code stays private to kin (dominant-lineage share 99.1%, N_{eff} 1.02): even under actively sustained diversity, and even when survival is made to depend on decoding non-kin, no public cross-lineage code bootstraps, a *pressure-vs-survival tension*. We also contribute a bidirectional self-audit that catches deflation as readily as inflation.

Submission type: **Late Breaking Abstracts**

Code & verdict artifacts: <https://github.com/yangchunxuan/grown-grounded-language> (archived at Zenodo, DOI 10.5281/zenodo.21074235); every result regenerates from committed, seeded verdict JSONs.

The question and the world

Where does a shared signalling code come from? Most computational accounts *train* a channel in (gradient descent on a task with a differentiable channel; Foerster et al., 2016; Lazaridou et al., 2017) or *design* one in (an imposed iterated-learning transmission bottleneck; Kirby et al., 2008; Ren et al., 2020). We take up the prior question: can a decodable code *grow* from random weights under **survival selection alone**, and if it does, is it a **public** code or a code private to **kin**?

The substrate is deliberately minimal and entirely hand-coded: a 48×48 lethal-vent world in which agents

metabolise, can die, and reproduce, and in which a single tied UnifiedIO code-table ($4 \times 5 \times 6$, $\mathcal{N}(0,1)$ init) is the only communicative apparatus. There is no gradient on communication, no oracle (agents never read world truth), and no borrowed weights. We measure mutual intelligibility (MII) behaviourally: route a referent through one agent's emit into another's decode and score exact-tuple recovery against a chance floor of $1/625 \approx 0.0016$. Within-founder (WF) and cross-founder (CF) MII separate a *private kin code* (WF high, CF at chance) from a *public* one (CF also above chance).

Co-evolution is real, but decomposes and stays kin-private

The control is not communication-ablated: both arms carry the full code-table and run the identical loop, and the only manipulated variable is the reproduction-fitness signal (survival vs. uniform-random, with the reproduction RNG byte-matched). This isolates selection from *clonal descent-convergence*. The random-fitness arm also fixates to $\sim$ one lineage, so its MII of **0.1174** ($\sim 73\times$ floor) is the descent baseline, not zero. At gen 28 the survival arm reaches **0.17219**, a paired win-margin of +0.0548, **95% CI** [0.0234, 0.0849] ($n=48$; Figure 1, panel a). Selection therefore *sharpens* a largely descent-determined channel rather than originating it.

A post-hoc probe ($n=24$; gen 28/56/100/150) shows this is a *transient acceleration*: the survival-arm MII holds at ~ 0.16 – 0.17 while the control's rises ($0.117 \rightarrow 0.138$) as clonal fixation accumulates decodability, so the paired margin decays from +0.0569 (CI excludes zero) to a CI including zero by gen 100. The headline is therefore scoped to 28 generations.

The code is *not public*. The survival arm fixates to a single founder lineage (dominant-lineage share **99.1%**, N_{eff} **1.02**, 47/48 seeds single-lineage), so WF runs ≈ 0.16 ($\sim 100\times$ chance) while CF sits at the floor. Almost every receiver an agent meets is kin, so a private code is selectively cheap (Hamilton, 1964).




Figure 1: **A channel grows, but only kin can read it.** (a) Co-evolution headline ($n=48$, gen 28): the survival arm beats the architecture-matched random-fitness control by a paired MII win-margin $+0.0548$ (95% CI $[0.0234, 0.0849]$); a post-hoc probe shows this is a transient acceleration. (b) Cross-founder MII stays at the chance floor (≈ 0.0016) across three cross-lineage interventions ($n=8$); where pressure bites, survival crashes. No public code bootstraps. Sources: committed verdict JSONs.

The wall: no public code bootstraps

Is kin-privacy merely an artifact of fast founder collapse? We sustained genuine multi-lineage coexistence by niching (tournament selection plus lineage fitness-sharing) at $\text{pop}=96$: the soft96 cell holds $N_{\text{eff}} 4.10$ over 20+ coexisting generations, yet CF stays at the floor (CF 0.00153 vs. floor 0.0016; the CF–floor CI straddles zero). Maintaining diversity is *necessary but not sufficient*.

Making survival itself depend on decoding non-kin does not help either. Across three heterogeneous interventions (Figure 1, panel b: C2X2 at 0.25 non-kin stays viable (alive 0.16) yet CF 0.00156; C2X3 at 0.5 non-kin alive 0.10, CF 0.00192; C2X at 100% non-kin crashes to alive 0.046), CF never rose above chance; where the pressure bit, it crashed survival. Non-kin codes are mutually undecodable, so forced cross-listening is acting on noise, which starves foraging: a *pressure-vs-survival tension*. All cross-lineage arms are $n=8$ (positive-or-inconclusive tier), so we report that a public code *did not* bootstrap under the regimes tested, never that none *can*.

A bidirectional self-audit

Each claim must survive an architecture-matched control, a fair baseline, ≥ 6 -seed bootstrap-CI win-margins, frozen pre-registration, a no-oracle redline, and code-grounded adversarial review. Its single methodological claim is *bidirectionality*: it catches deflation as readily as inflation. Two in-scope catches show this: a pilot/formal false *negative* (a formal arm compared against a pilot control had buried a real effect, with generation count the load-bearing knob), and an RNG-offset bug in the headline control whose fix left the conclusion intact and slightly *strengthened* the margin ($+0.045 \rightarrow +0.0548$).

Contribution and scope

This separates three claims often conflated in emergent-language work: a decodable code can *evolve*, it can *matter* for survival, and it can still *fail to become public*. A decodable code grows from scratch under survival selection

but stays kin-private; selection *accelerates* a clonal-descent channel rather than originating a public one; and no public code bootstraps even under sustained diversity and direct cross-lineage pressure. The results are scoped to one hand-coded world, one architecture, and per-stage timescales (28–48 generations). We read the wall as motivating a richer world and a consequence-driven learner rather than more selection pressure. All code and seeded verdict JSONs are committed and reproducible.

Acknowledgements

This work received no specific funding and the author declares no competing interests. AI-based tools were used to assist with drafting and language editing; the author verified all content and takes full responsibility for the manuscript.

References

- Foerster, J., et al. (2016). Learning to communicate with deep multi-agent reinforcement learning. *NeurIPS*.
- Hamilton, W. D. (1964). The genetical evolution of social behaviour. *J. Theor. Biol.*, 7(1).
- Kirby, S., Cornish, H., and Smith, K. (2008). Cumulative cultural evolution in the laboratory. *PNAS*, 105(31).
- Lazaridou, A., Peysakhovich, A., and Baroni, M. (2017). Multi-agent cooperation and the emergence of (natural) language. *ICLR*.
- Ren, Y., et al. (2020). Compositional languages emerge in a neural iterated learning model. *ICLR*.